

1. Project Overview

- A **multi-agent system** for querying PDFs, web, and Arxiv research papers.
 - Integrates a **Decision Agent** to intelligently route queries and an **Answer Synthesizer Agent** to produce a clean, human-readable answer.
 - Minimal frontend for PDF upload, query input, and result display.
-

2. Architecture

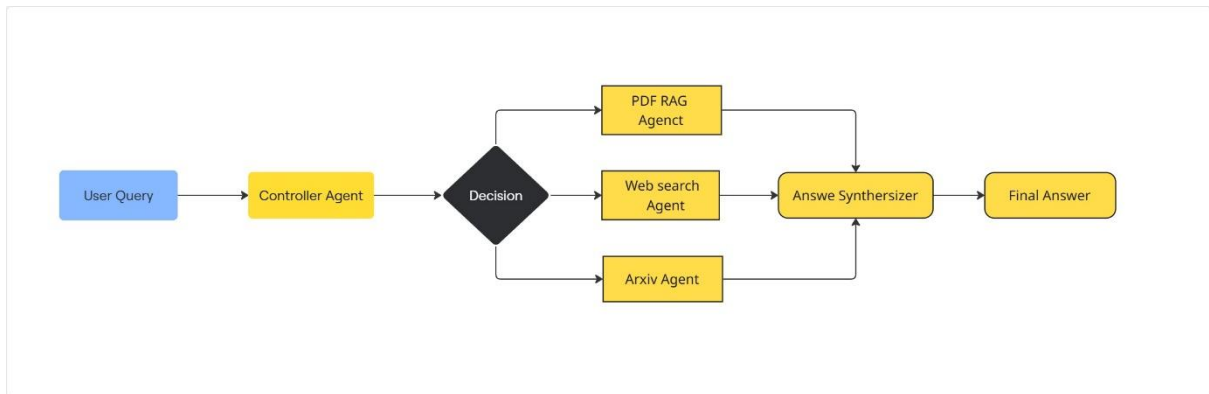
Agents:

1. **PDF RAG Agent (pdf_rag.py)**
 - Handles uploaded PDFs.
 - Uses RAG to retrieve relevant chunks of text.
2. **Web Search Agent (web_search.py)**
 - Handles queries about latest news or events.
 - Fetches snippets from web sources.
3. **Arxiv Agent (arxiv_agent.py)**
 - Handles research-oriented queries.
 - Fetches summaries from Arxiv papers.
4. **Decision Agent (decision_agent.py)**
 - Decides which agents to call based on the query and available files.
 - Returns JSON: {"agents": [...], "rationale": "..."}.
5. **Answer Synthesizer (answer_synthesizer.py)**
 - Combines snippets from multiple agents.
 - Produces a **clear, concise, factually correct** answer.

Controller (controller.py)

- Central orchestrator.
- Receives the query, asks the Decision Agent which agents to call.
- Collects results from each agent.
- Calls Answer Synthesizer for final output.

User Flow Chart



3. Controller Decision Logic

- **Checks for:**
 - File uploaded → routes to PDF agent.
 - Research keywords → routes to Arxiv agent.
 - Latest news / update keywords → routes to Web agent.
 - **Default: routes to Web agent if no other matches.**
-

4. Safety & Privacy

- API keys stored in .env (never pushed to GitHub).
 - No user data is logged outside the system.
 - PDFs are processed locally; no external storage except temporary snippets.
-

5. Limitations

- PDFs with heavy images require Tesseract OCR, which is not installed on all deployment platforms.
 - Gemini API usage may incur cost limits.
 - Web scraping is basic; may not handle all websites.
 - Arxiv summaries rely on available metadata and abstracts.
-

6. Sample Usage

1. Upload a PDF via frontend.
2. Ask a query like "Minimize this PDF" or "Latest research about LLM engineering".

3. System routes to appropriate agents and returns a synthesized answer.
4. Logs show which agents were used and the rationale.

6. Folder Structure / Deliverables

```
backend/  
├─ app/  
│   └─ agent/  
│       ├── pdf_rag.py  
│       ├── web_search.py  
│       ├── arxiv_agent.py  
│       ├── decision_agent.py  
│       ├── answer_synthesizer.py  
│       └─ controller.py  
│   └─ utils/  
│       └─ pdf_utils.py  
│   └─ main.py  
└─ uploads/  
frontend/  
├─ index.html  
└─ app.js
```

8. Conclusion

- Modular, easily extendable multi-agent system.
- Demonstrates RAG, web search, and decision-making orchestration.
- Deployed in Render.